

## Assignment No. 1

```
# importing libraries
import pandas as pd
import numpy as np
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# Loading dataset
df = pd.read_csv("covid_19_clean_complete_2022.csv")
df.drop('Province/State', axis=1, inplace=True)
df.head()

# data preprocessing - detecting NaN values and using describe() function
df.isna().any()
df.describe()

# shape of dataset (dimensions)
df.shape

# data formatting and normalization
df.dtypes

df['Date'] = pd.to_datetime(df['Date'])
df['Country/Region'] = df['Country/Region'].astype('string')
df.dtypes

# handling categorical values
# dropping the categorical variable column
# df['new_col'] = df['some_col'].map({'value_1': 1, 'value_2': 2})
df = df.drop(['WHO Region', 'Country/Region'], axis=1)
df.head()
```

## Output

Country/Region	False
Lat	True
Long	True
Date	False
Confirmed	False
Deaths	False
Recovered	False
Active	False

WHO Region        True  
dtype: bool

	Lat	Long	Confirmed	Deaths	Recovered	Active
<b>count</b>	213348.000000	213348.000000	2.148940e+05	214894.000000	2.148940e+05	2.148940e+05
<b>mean</b>	20.528131	22.735337	4.578132e+05	9310.764693	1.079987e+05	3.405037e+05
<b>std</b>	25.899139	76.304185	2.708770e+06	47497.835275	8.470111e+05	2.516382e+06
<b>min</b>	-71.949900	-178.116500	0.000000e+00	0.000000	0.000000e+00	-1.638280e+05
<b>25%</b>	6.426991	-27.932425	2.530000e+02	2.000000	0.000000e+00	1.600000e+01
<b>50%</b>	22.233350	21.752000	5.223000e+03	71.000000	4.500000e+01	1.243000e+03
<b>75%</b>	41.166070	88.658375	9.892275e+04	1675.000000	5.115750e+03	2.644675e+04
<b>max</b>	71.706900	178.065000	7.925051e+07	958144.000000	3.097475e+07	7.829236e+07

Country/Region    object

Lat                float64

Long               float64

Date                object

Confirmed          int64

Deaths             int64

Recovered          int64

Active             int64

WHO Region        object

dtype: object

## Assignment No. 2

```
# Commented out IPython magic to ensure Python compatibility.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
# %matplotlib inline
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

df = pd.read_csv("Academic-Performance-Dataset.csv")
df

df.shape

df.dtypes

df.isna().sum()
cols_with_na = []
for col in df.columns:
    if df[col].isna().any():
        cols_with_na.append(col)

cols_with_na
for col in cols_with_na:
    col_dt = df[col].dtypes
    if (col_dt == 'int64' or col_dt == 'float64'):
        outliers = (df[col] < 0) | (100 < df[col])
        df.loc[outliers, col] = np.nan
        df[col] = df[col].fillna(df[col].mean())
    else:
        df[col] = df[col].fillna(method='ffill')

df
df['Total
Marks']=df['Phy_marks']+df['Che_marks']+df['EM1_marks']+df['PPS_marks']+df['SME_m
arks']
df['Percentage']=df['Total Marks']/5

df
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (9, 6)
df_list = ['Attendance', 'Phy_marks', 'Che_marks', 'EM1_marks', 'PPS_marks',
'SME_marks']
```

```

fig, axes = plt.subplots(2, 3)
fig.set_dpi(120)

count=0
for r in range(2):
    for c in range(3):
        _ = df[df_list[count]].plot(kind = 'box', ax=axes[r,c])
        count+=1
Q1 = df['Che_marks'].quantile(0.25)
Q3 = df['Che_marks'].quantile(0.75)
IQR = Q3 - Q1

Lower_limit = Q1 - 1.5 * IQR
Upper_limit = Q3 + 1.5 * IQR

print(f'Q1 = {Q1}, Q3 = {Q3}, IQR = {IQR}, Lower_limit = {Lower_limit},
Upper_limit = {Upper_limit}')

df[(df['Che_marks'] < Lower_limit) | (df['Che_marks'] > Upper_limit)]
def BinningFunction(column, cut_points, labels = None) :
    break_points=[column.min()] + cut_points + [column.max( )]
    print('Gradding According to percentage \n>60 = F \n60-70 = B \n70-80 =
A\n80-100 = O')
    return pd.cut(column, bins=break_points, labels=labels, include_lowest=True)

cut_points=[60, 70, 80]
labels=['F', 'B', 'A', 'O']
df['Grade']=BinningFunction(df['Percentage'], cut_points, labels)

df

```

## Output

```


(20, 12)
Rollno      int64
Name        object
Gender      object
Branch      object

```

Attendance float64  
Phy\_marks float64  
Che\_marks float64  
EM1\_marks float64  
PPS\_marks float64  
SME\_marks float64  
Total Marks int64  
Percentage float64  
dtype: object

Rollno 0  
Name 2  
Gender 0  
Branch 0  
Attendance 0  
Phy\_marks 1  
Che\_marks 3  
EM1\_marks 2  
PPS\_marks 1  
SME\_marks 0  
Total Marks 0  
Percentage 0  
dtype: int64

```
['Name', 'Phy_marks', 'Che_marks', 'EM1_marks', 'PPS_marks']
```



Q1 = 76.0, Q3 = 86.25, IQR = 10.25, Lower\_limit = 60.625, Upper\_limit = 101.625

Rollno	Name	Gender	Branch	Attendance	Phy_marks	Che_marks	EM1_marks				
	PPS_marks		SME_marks	Total Marks	Percentage						
7	8	Ishaan	M	ENTC	75.0	66.0	51.0	83.0	69.611111	76.0	
	345.611111		69.122222								
14	15	Maryam		F	IT	64.0	87.0	60.0	90.0	65.000000	90.0
	392.000000		78.400000								

Gradding According to percentage

>60 = F  
 60-70 = B  
 70-80 = A  
 80-100 = O

### Assignment No. 3

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

df = pd.read_csv("iris.csv")
df.head()

'Iris-setosa'
setosa = df['Species'] == 'Iris-setosa'
df[setosa].describe()
'Iris-versicolor'
versicolor = df['Species'] == 'Iris-versicolor'
df[versicolor].describe()
'Iris-virginica'
virginica = df['Species'] == 'Iris-virginica'
df[virginica].describe()

df.dtypes
df.dtypes.value_counts()
```

#### Output

##### Iris-setosa

Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	
count	50.00000	50.00000	50.000000	50.000000	50.00000
mean	25.50000	5.00600	3.418000	1.464000	0.24400
std	14.57738	0.35249	0.381024	0.173511	0.10721
min	1.00000	4.30000	2.300000	1.000000	0.10000
25%	13.25000	4.80000	3.125000	1.400000	0.20000
50%	25.50000	5.00000	3.400000	1.500000	0.20000
75%	37.75000	5.20000	3.675000	1.575000	0.30000
max	50.00000	5.80000	4.400000	1.900000	0.60000

##### Iris-versicolor

Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
----	---------------	--------------	---------------	--------------

count	50.00000	50.000000	50.000000	50.000000	50.000000
mean	75.50000	5.936000	2.770000	4.260000	1.326000
std	14.57738	0.516171	0.313798	0.469911	0.197753
min	51.00000	4.900000	2.000000	3.000000	1.000000
25%	63.25000	5.600000	2.525000	4.000000	1.200000
50%	75.50000	5.900000	2.800000	4.350000	1.300000
75%	87.75000	6.300000	3.000000	4.600000	1.500000
max	100.00000	7.000000	3.400000	5.100000	1.800000

Iris-virginica

Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	50.00000	50.00000	50.000000	50.000000
mean	125.50000	6.58800	2.974000	5.552000
std	14.57738	0.63588	0.322497	0.551895
min	101.00000	4.90000	2.200000	4.500000
25%	113.25000	6.22500	2.800000	5.100000
50%	125.50000	6.50000	3.000000	5.550000
75%	137.75000	6.90000	3.175000	5.875000
max	150.00000	7.90000	3.800000	6.900000

```

Id          int64
SepalLengthCm  float64
SepalWidthCm  float64
PetalLengthCm  float64
PetalWidthCm  float64
Species       object
dtype: object
float64  4
int64    1
object   1
dtype: int64

```



## Assignment No. 4

```
# Commented out IPython magic to ensure Python compatibility.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
# %matplotlib inline

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# Importing DataSet and take a look at Data
Boston = pd.read_csv("boston.csv")
Boston.head()

Boston.info()
Boston.describe()

Boston.plot.scatter('RM', 'MEDV', figsize=(6, 6));

plt.subplots(figsize=(10,8))
sns.heatmap(Boston.corr(), cmap = 'coolwarm', annot = True, fmt = '.1f');

X = Boston[Boston.columns[:-1]]
Y = Boston['MEDV']

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler



# Split DataSet
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3)
sc_X = StandardScaler()
X_train_ = sc_X.fit_transform(X_train)
X_test_ = sc_X.transform(X_test)


print(f'Train Dataset Size - X: {X_train.shape}, Y: {Y_train.shape}')
print(f'Test Dataset Size - X: {X_test.shape}, Y: {Y_test.shape}')

# Model Building
lm = LinearRegression()
lm.fit(X_train_, Y_train)
predictions = lm.predict(X_test_)
```



std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
	7.141062	9.197104				
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
	1.730000	5.000000				
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
	6.950000	17.025000				
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
	11.360000	21.200000				
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
	16.955000	25.000000				
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000
	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000
	37.970000	50.000000				





## Assignment No. 5

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

df = pd.read_csv('Social_Network_Ads.csv')
df.head()
df.info()
df.describe()

X = df[['Age', 'EstimatedSalary']]
Y = df['Purchased']

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.25,
random_state = 0)
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)

print(f'Train Dataset Size - X: {X_train.shape}, Y: {Y_train.shape}')
print(f'Test Dataset Size - X: {X_test.shape}, Y: {Y_test.shape}')

from sklearn.linear_model import LogisticRegression

lm = LogisticRegression(random_state = 0, solver='lbfgs' )
lm.fit(X_train, Y_train)
predictions = lm.predict(X_test)

plt.figure(figsize=(6, 6));
sns.regplot(x = X_test[:, 1], y = predictions, scatter_kws={'s':5});
plt.scatter(X_test[:, 1], Y_test, marker = '+');
plt.xlabel("User's Estimated Salary");
plt.ylabel('Ads Purchased');
plt.title('Regression Line Tracing');
```

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

cm = confusion_matrix(Y_test, predictions)
print(f'''Confusion matrix :\n
          | Positive Prediction\t| Negative Prediction
-----+-----+-----
Positive Class | True Positive (TP) {cm[0, 0]}\t| False Negative (FN) {cm[0, 1]}
-----+-----+-----
Negative Class | False Positive (FP) {cm[1, 0]}\t| True Negative (TN) {cm[1, 1]}\n\n''')

cm = classification_report(Y_test, predictions)
print('Classification report : \n', cm)

# Visualizing the Training set results
from matplotlib.colors import ListedColormap
X_set, y_set = X_train, Y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))

plt.figure(figsize=(9, 7.5));
plt.contourf(X1, X2, lm.predict(np.array([X1.ravel(),
X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.6, cmap = ListedColormap(('red', 'green')));
plt.xlim(X1.min(), X1.max());
plt.ylim(X2.min(), X2.max());
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                color = ListedColormap(('red', 'green'))(i), label = j);
plt.title('Logistic Regression (Training set)');
plt.xlabel('Age');
plt.ylabel('Estimated Salary');
plt.legend();
plt.show();

# Visualizing the Test set results
from matplotlib.colors import ListedColormap
X_set, y_set = X_test, Y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))

```

```
plt.figure(figsize=(9, 7.5));
plt.contourf(X1, X2, lm.predict(np.array([X1.ravel(),
X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.6, cmap = ListedColormap(('red', 'green')));
plt.xlim(X1.min(), X1.max());
plt.ylim(X2.min(), X2.max());
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
               color = ListedColormap(('red', 'green'))(i), label = j);
plt.title('Logistic Regression (Test set)');
plt.xlabel('Age');
plt.ylabel('Estimated Salary');
plt.legend();
plt.show();
```

## Output

User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19.0	19000.0 0
1	15810944	Male	35.0	20000.0 0
2	15668575	Female	26.0	43000.0 0
3	15603246	Female	27.0	57000.0 0
4	15804002	Male	19.0	76000.0 0

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 400 entries, 0 to 399

Data columns (total 5 columns):


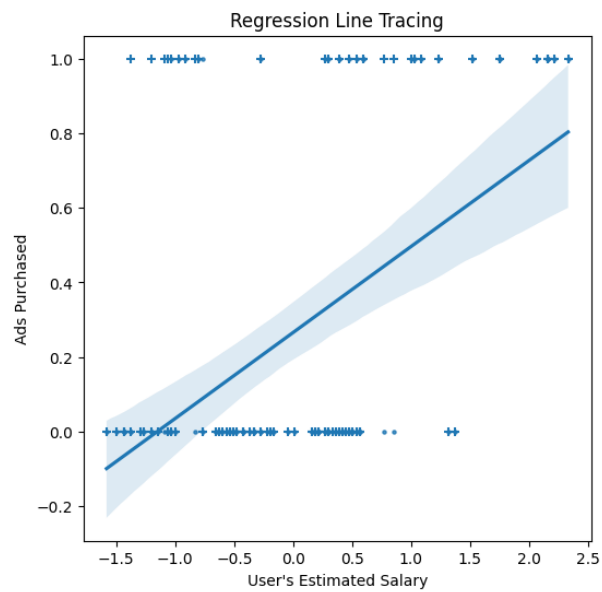
#	Column	Non-Null Count	Dtype
0	User ID	400 non-null	int64
1	Gender	400 non-null	object
2	Age	400 non-null	float64
3	EstimatedSalary	400 non-null	float64
4	Purchased	400 non-null	int64

dtypes: float64(2), int64(2), object(1)

memory usage: 15.8+ KB

User ID	Age	EstimatedSalary	Purchased
---------	-----	-----------------	-----------

count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000





## Assignment No. 6

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

df = pd.read_csv('iris.csv')
df.head()

X = df.iloc[:, :4].values
Y = df['Species'].values

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
random_state = 0)
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)

print(f'Train Dataset Size - X: {X_train.shape}, Y: {Y_train.shape}')
print(f'Test Dataset Size - X: {X_test.shape}, Y: {Y_test.shape}')

from sklearn.naive_bayes import GaussianNB

classifier = GaussianNB()
classifier.fit(X_train, Y_train)
predictions = classifier.predict(X_test)

mapper = {'setosa': 0, 'versicolor': 1, 'virginica': 2}
predictions_ = [mapper[i] for i in predictions]

fig, axs = plt.subplots(2, 2, figsize = (12, 10), constrained_layout = True);
_ = fig.suptitle('Regression Line Tracing')

for i in range(4):
    x, y = i // 2, i % 2
    _ = sns.regplot(x = X_test[:, i], y = predictions_, ax=axs[x, y])
    _ = axs[x, y].scatter(X_test[:, i][::-1], Y_test[::-1], marker = '+',
color="white")
```

```

_ = axs[x, y].set_xlabel(df.columns[i + 1][: -2])

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report


cm = confusion_matrix(Y_test, predictions)
print(f'''Confusion matrix :\n
          | Positive Prediction\t| Negative Prediction
-----+-----+-----
Positive Class | True Positive (TP) {cm[0, 0]}\t| False Negative (FN) {cm[0, 1]}
-----+-----+-----
Negative Class | False Positive (FP) {cm[1, 0]}\t| True Negative (TN) {cm[1, 1]}\n\n''')

cm = classification_report(Y_test, predictions)
print('Classification report : \n', cm)

```

## Output

Id		SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	setosa
1	2	4.9	3.0	1.4	0.2	setosa
2	3	4.7	3.2	1.3	0.2	setosa
3	4	4.6	3.1	1.5	0.2	setosa
4	5	5.0	3.6	1.4	0.2	setosa



## Assignment No. 7

```
# import required module
from sklearn.feature_extraction.text import TfidfVectorizer

# assign documents
d0 = 'Just Code It'
d1 = 'Code'
d2 = 'Programming'

# merge documents into a single corpus
string = [d0, d1, d2]

# create object
tfidf = TfidfVectorizer()

# get tf-df values
result = tfidf.fit_transform(string)

# get idf values
print('\nidf values:')
for ele1, ele2 in zip(tfidf.get_feature_names(), tfidf.idf_):
    print(ele1, ': ', ele2)

# get indexing
print('\nWord indexes:')
print(tfidf.vocabulary_)

# display tf-idf values
print('\ntf-idf value:')
print(result)

# in matrix form
print('\ntf-idf values in matrix form:')
print(result.toarray())
```

Output:-

Word indexes:

{'just': 2, 'code': 0, 'it': 1, 'programming': 3}

tf-idf value:

(0, 1) 0.6227660078332259

(0, 0) 0.4736296010332684

(0, 2) 0.6227660078332259

(1, 0) 1.0

(2, 3) 1.0

tf-idf values in matrix form:

[[0.4736296 0.62276601 0.62276601 0. ]

[1. 0. 0. 0. ]

[0. 0. 0. 1. ]]

## Assignment No. 8

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns


dataset = sns.load_dataset('titanic')

dataset.head()

sns.histplot(dataset['fare'], kde=True, linewidth=0);

sns.jointplot(x='age', y='fare', data=dataset);
```

Output:



## Assignment No. 9

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns


dataset = sns.load_dataset('titanic')

dataset.head()

sns.boxplot(x='sex', y='age', data=dataset, hue="survived");
```

### Output

survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	
	embark_town	alive	alone									
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN
	Southampton	no	False									
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C
	Cherbourg	yes	False									
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN
	Southampton	yes	True									
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C
	Southampton	yes	False									
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN
	Southampton	no	True									



## Assignment No. 10

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

df = pd.read_csv('iris.csv')
df.head()

df.info()

np.unique(df["Species"])

df.describe()

import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt

fig, axes = plt.subplots(2, 2, figsize=(12, 6), constrained_layout = True)

for i in range(4):
    x, y = i // 2, i % 2
    _ = axes[x, y].hist(df[df.columns[i + 1]])
    _ = axes[x, y].set_title(f"Distribution of {df.columns[i + 1][: -2]}")

data_to_plot = df[df.columns[1:-1]]

fig, axes = plt.subplots(1, figsize=(12,8))
bp = axes.boxplot(data_to_plot)
```

## Output

Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species	
0	1	5.1	3.5	1.4	0.2	setosa
1	2	4.9	3.0	1.4	0.2	setosa
2	3	4.7	3.2	1.3	0.2	setosa
3	4	4.6	3.1	1.5	0.2	setosa
4	5	5.0	3.6	1.4	0.2	setosa

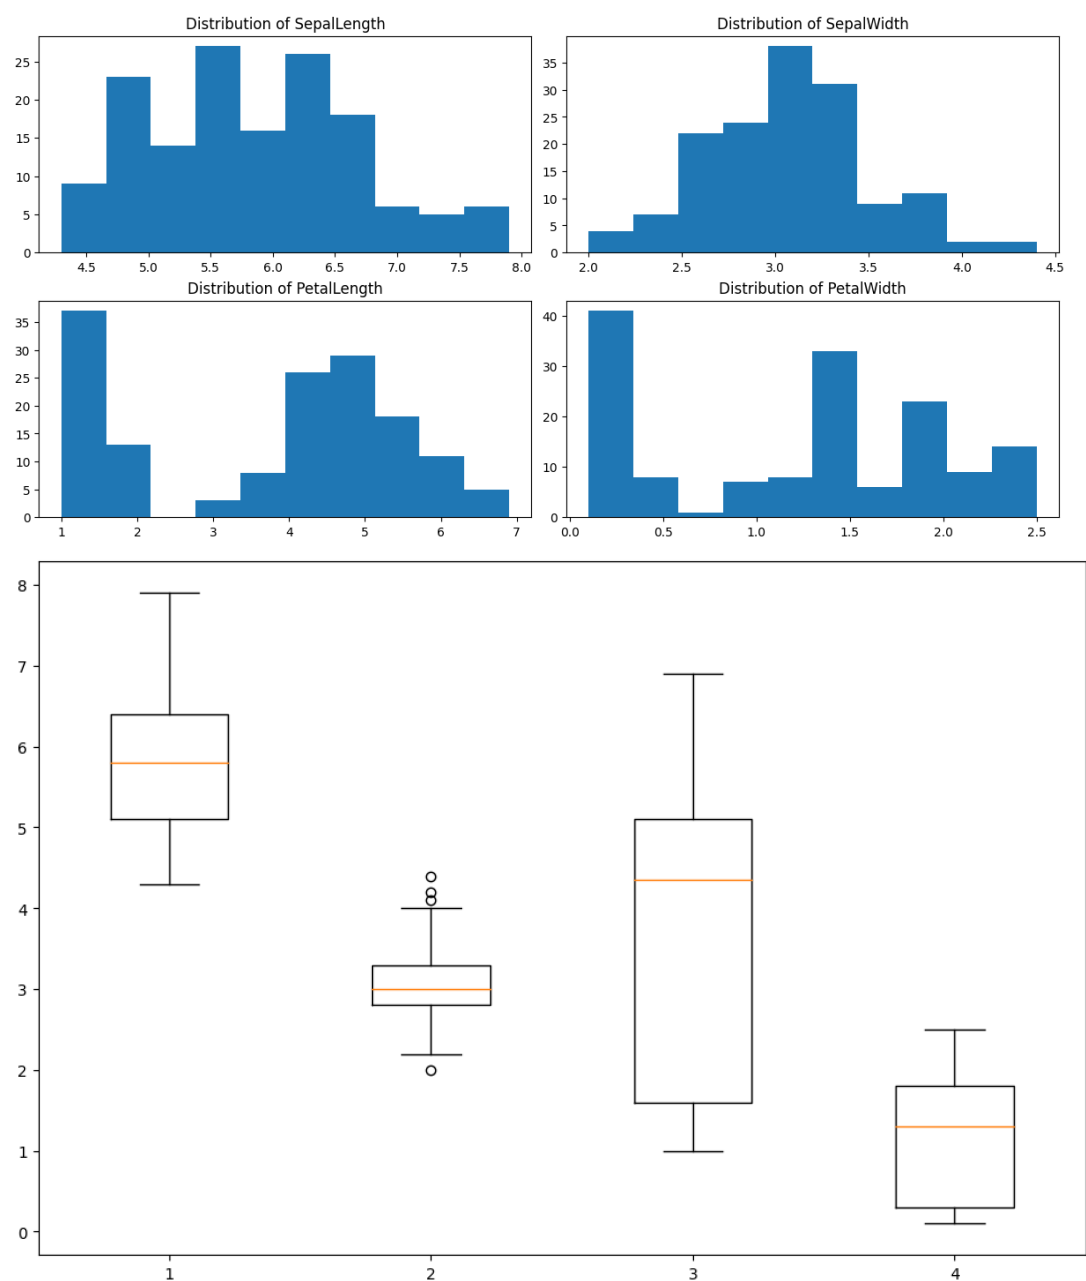
RangeIndex: 150 entries, 0 to 149

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	Id	150 non-null	int64
1	SepalLengthCm	150 non-null	float64
2	SepalWidthCm	150 non-null	float64
3	PetalLengthCm	150 non-null	float64
4	PetalWidthCm	150 non-null	float64
5	Species	150 non-null	object

dtypes: float64(4), int64(1), object(1)

memory usage: 7.2+ KB





```
// WC_Runner.java

package com.wc;

import java.io.IOException;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.FileInputFormat;

import org.apache.hadoop.mapred.FileOutputFormat;

import org.apache.hadoop.mapred.JobClient;

import org.apache.hadoop.mapred.JobConf;

import org.apache.hadoop.mapred.TextInputFormat;

import org.apache.hadoop.mapred.TextOutputFormat;

public class WC_Runner {

    public static void main(String[] args) throws IOException {

        JobConf conf = new JobConf(WC_Runner.class);

        conf.setJobName("WordCount");

        conf.setOutputKeyClass(Text.class);

        conf.setOutputValueClass(IntWritable.class);

        conf.setMapperClass(WC_Mapper.class);

        conf.setCombinerClass(WC_Reducer.class);

        conf.setReducerClass(WC_Reducer.class);

        conf.setInputFormat(TextInputFormat.class);

        conf.setOutputFormat(TextOutputFormat.class);

        FileInputFormat.setInputPaths(conf,new Path(args[0]));

        FileOutputFormat.setOutputPath(conf,new Path(args[1]));

        JobClient.runJob(conf);

    }

}
```

Assignment No. 11

// WC\_Mapper.java

package com.wc;

import java.io.IOException;

import java.util.StringTokenizer;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.LongWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.MapReduceBase;

import org.apache.hadoop.mapred.Mapper;

import org.apache.hadoop.mapred.OutputCollector;

import org.apache.hadoop.mapred.Reporter;

public class WC\_Mapper extends MapReduceBase implements

Mapper<LongWritable,Text,Text,IntWritable>{

private final static IntWritable one = new IntWritable(1);

private Text word = new Text();

public void map(

LongWritable key,

Text value,

OutputCollector<Text,IntWritable> output,

Reporter reporter

) throws IOException {

String line = value.toString();

StringTokenizer tokenizer = new StringTokenizer(line);

while (tokenizer.hasMoreTokens()){

word.set(tokenizer.nextToken());

output.collect(word, one);

}

```

}
}
// WC_Reducer.java

package com.wc;

import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WC_Reducer extends MapReduceBase implements
Reducer<Text,IntWritable,Text,IntWritable> {

    public void reduce(
        Text key,
        Iterator<IntWritable> values,
        OutputCollector<Text,IntWritable> output,
        Reporter reporter
    ) throws IOException {
        int sum=0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key,new IntWritable(sum));
    }
}

```

Input:

HDFS is a storage unit of Hadoop

MapReduce is a processing tool for Hadoop

Output:

HDFS 1

Hadoop 2

MapReduce 1

a 2

for 1

is 2

of 1

processing 1

storage 1

tool 1

unit 1

## Assignment No. 12

```
// SalesCountryRunner.java

package SalesCountry;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class SalesCountryRunner {

    public static void main(String[] args) {
        JobClient my_client = new JobClient();

        // Create a configuration object for the job
        JobConf job_conf = new JobConf(SalesCountryDriver.class);

        // Set a name of the Job
        job_conf.setJobName("SalePerCountry");

        // Specify data type of output key and value
        job_conf.setOutputKeyClass(Text.class);
        job_conf.setOutputValueClass(IntWritable.class);

        // Specify names of Mapper and Reducer Class
        job_conf.setMapperClass(SalesCountry.SalesMapper.class);
        job_conf.setReducerClass(SalesCountry.SalesCountryReducer.class);

        // Specify formats of the data type of Input and output
        job_conf.setInputFormat(TextInputFormat.class);
        job_conf.setOutputFormat(TextOutputFormat.class);

        // Set input and output directories using command line arguments,
        //arg[0] = name of input directory on HDFS, and arg[1] = name of output
        //directory to be created to store the output file.

        FileInputFormat.setInputPaths(job_conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(job_conf, new Path(args[1]));
    }
}
```

```

my_client.setConf(job_conf);
try {
// Run the job
JobClient.runJob(job_conf);
} catch (Exception e) {
e.printStackTrace();
}
}
}

// SalesMapper.java
package SalesCountry;
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class SalesMapper extends MapReduceBase implements Mapper<LongWritable, Text,
Text,
IntWritable> {
private final static IntWritable one = new IntWritable(1);
public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
output, Reporter reporter) throws IOException {
String valueString = value.toString();
String[] SingleCountryData = valueString.split(",");
output.collect(new Text(SingleCountryData[7]), one);
}
}

// SalesCountryReducer.java
package SalesCountry;

```

```

import java.io.IOException;
import java.util.*;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class SalesCountryReducer extends MapReduceBase implements Reducer<Text,
IntWritable,
Text, IntWritable> {
public void reduce(Text t_key, Iterator<IntWritable> values,
OutputCollector<Text,IntWritable> output, Reporter reporter) throws IOException {
Text key = t_key;
int frequencyForCountry = 0;
while (values.hasNext()) {
// replace type of value with the actual type of our value
IntWritable value = (IntWritable) values.next();
frequencyForCountry += value.get();
}
output.collect(key, new IntWritable(frequencyForCountry));
}
}

```

Output:

```

Argentina 1
Australia 38
Austria 7
Bahrain 1
Belgium 8
Bermuda 1
Brazil 5
Bulgaria 1

```

CO 1

Canada 76

Cayman Isls 1

China 1

Costa Rica 1

Country 1

Czech Republic 3

Denmark 15

Dominican Republic 1

Finland 2

France 27

Germany 25

Greece 1

Guatemala 1

Hong Kong 1

Hungary 3

Iceland 1

India 2

Ireland 49

Israel 1

Italy 15

Japan 2

Jersey 1

Kuwait 1

Latvia 1

Luxembourg 1

Malaysia 1

Malta 2



Mauritius 1  
Moldova 1  
Monaco 2  
Netherlands 22  
New Zealand 6  
Norway 16  
Philippines 2  
Poland 2  
Romania 1  
Russia 1  
South Africa 5  
South Korea 1  
Spain 12  
Sweden 13  
Switzerland 36  
Thailand 2  
The Bahamas 2  
Turkey 6  
Ukraine 1  
United Arab Emirates 6  
United Kingdom 100  
United States 462

## Assignment No. 13

```
// MaxTemperatureDriver.java

package MaxMinTemp;

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class MaxTemperatureDriver extends Configured implements Tool{

    public int run(String[] args) throws Exception {

        if(args.length !=2) {

            System.err.println("Usage: MaxTemperatureDriver <input path>
<outputpath>");

            System.exit(-1);

        }

        Job job = new Job();

        job.setJarByClass(MaxTemperatureDriver.class);

        job.setJobName("Max Temperature");

        FileInputFormat.addInputPath(job, new Path(args[0]));

        FileOutputFormat.setOutputPath(job,new Path(args[1]));

        job.setMapperClass(MaxTemperatureMapper.class);

        job.setReducerClass(MaxTemperatureReducer.class);

        job.setOutputKeyClass(Text.class);
```

```

job.setOutputValueClass(IntWritable.class);
System.exit(job.waitForCompletion(true) ? 0:1);
boolean success = job.waitForCompletion(true);
return success ? 0 : 1;
}

public static void main(String[] args) throws Exception {
    MaxTemperatureDriver driver = new MaxTemperatureDriver();
    int exitCode = ToolRunner.run(driver, args);
    System.exit(exitCode);
}
}

// MaxTemperatureMapper.java
package MaxMinTemp;
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class MaxTemperatureMapper extends Mapper<LongWritable, Text, Text,
IntWritable> {
    private static final int MISSING = 9999;

    @Override
    public void map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException {
        String line = value.toString();
        String year = line.substring(15, 19);
        int airTemperature;
        if (line.charAt(87) == '+') { // parseInt doesn't like leading plus signs

```

```

airTemperature = Integer.parseInt(line.substring(88, 92));
} else {
airTemperature = Integer.parseInt(line.substring(87, 92));
}
String quality = line.substring(92, 93);
if (airTemperature != MISSING && quality.matches("[01459]")) {
context.write(new Text(year), new IntWritable(airTemperature));
}
}
}

// MaxTemperatureReducer.java
package MaxMinTemp;
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class MaxTemperatureReducer extends Reducer<Text, IntWritable, Text,
IntWritable> {
@Override
public void reduce(Text key, Iterable<IntWritable> values, Context context)
throws IOException, InterruptedException {
int maxVal = Integer.MIN_VALUE;
for (IntWritable value : values) {
maxVal = Math.max(maxVal, value.get());
}
context.write(key, new IntWritable(maxVal));
}
}

```

OUTPUT:

1901 317

1902 244

1903 289

1904 256

1905 283

1906 294

1907 283

1908 289

1909 278

1910 294

1911 306

1912 322

1913 300

1914 333

1915 294

1916 278

1917 317

1918 322

1919 378

1920 294